

In [2]:

```
# You can import data from various sources into your Pandas
# dataframe.
# A CSV file is a type of file where each line contains a single
# record, and all the columns are separated from each other via
# a comma.
# You can read CSV files using the read_csv() function of the
# Pandas dataframe, as shown below.
```

```
import pandas as pd
titanic_data = pd.read_csv("titanic.csv")
titanic_data.head()
```

```
# If you print the dataframe header, you should see that the
# header contains first five rows
```

Out[2]:

|   | survived | pclass | sex | age    | sibsp | parch | fare | embarked | class | who   | adult_male | deck  |
|---|----------|--------|-----|--------|-------|-------|------|----------|-------|-------|------------|-------|
| 0 | NaN      | 0      | 3   | male   | 22.0  | 1     | 0    | 7.2500   | S     | Third | man        | True  |
| 1 | NaN      | 1      | 1   | female | 38.0  | 1     | 0    | 71.2833  | C     | First | woman      | False |
| 2 | NaN      | 1      | 3   | female | 26.0  | 0     | 0    | 7.9250   | S     | Third | woman      | False |
| 3 | NaN      | 1      | 1   | female | 35.0  | 1     | 0    | 53.1000  | S     | First | woman      | False |
| 4 | NaN      | 0      | 3   | male   | 35.0  | 0     | 0    | 8.0500   | S     | Third | man        | True  |

In [3]:

```
import pandas as pd
titanic_data = pd.read_csv("titanic.csv")
titanic_data.tail()
```

```
# If you print the dataframe tail, you should see that the
# tail contains last five rows
```

Out[3]:

|     | survived | pclass | sex | age    | sibsp | parch | fare | embarked | class | who    | adult_male | deck |
|-----|----------|--------|-----|--------|-------|-------|------|----------|-------|--------|------------|------|
| 886 | NaN      | 0      | 2   | male   | 27.0  | 0     | 0    | 13.00    | S     | Second | man        |      |
| 887 | NaN      | 1      | 1   | female | 19.0  | 0     | 0    | 30.00    | S     | First  | woman      | F    |
| 888 | NaN      | 0      | 3   | female | NaN   | 1     | 2    | 23.45    | S     | Third  | woman      | F    |
| 889 | NaN      | 1      | 1   | male   | 26.0  | 0     | 0    | 30.00    | C     | First  | man        |      |
| 890 | NaN      | 0      | 3   | male   | 32.0  | 0     | 0    | 7.75     | Q     | Third  | man        |      |

In [15]:

```
# To handle missing numerical data, we can use statistical  
# techniques. The use of statistical techniques or algorithms to  
# replace missing values with statistically generated values is  
# called imputation.
```

```
import matplotlib.pyplot as plt  
import seaborn as sns  
plt.rcParams["figure.figsize"] = [8,6]  
sns.set_style("darkgrid")  
titanic_data = sns.load_dataset('titanic')  
titanic_data.head()
```

Out[15]:

|   | survived | pclass | sex    | age  | sibsp | parch | fare    | embarked | class | who   | adult_male |
|---|----------|--------|--------|------|-------|-------|---------|----------|-------|-------|------------|
| 0 | 0        | 3      | male   | 22.0 | 1     | 0     | 7.2500  | S        | Third | man   | True       |
| 1 | 1        | 1      | female | 38.0 | 1     | 0     | 71.2833 | C        | First | woman | False      |
| 2 | 1        | 3      | female | 26.0 | 0     | 0     | 7.9250  | S        | Third | woman | False      |
| 3 | 1        | 1      | female | 35.0 | 1     | 0     | 53.1000 | S        | First | woman | False      |
| 4 | 0        | 3      | male   | 35.0 | 0     | 0     | 8.0500  | S        | Third | man   | True       |

In [16]:

```
# Let's filter some of the numeric columns from the dataset and  
# see if they contain any missing values.
```

```
titanic_data = titanic_data[["survived", "pclass", "age", "fare"]]  
titanic_data.head()
```

Out[16]:

|   | survived | pclass | age  | fare    |
|---|----------|--------|------|---------|
| 0 | 0        | 3      | 22.0 | 7.2500  |
| 1 | 1        | 1      | 38.0 | 71.2833 |
| 2 | 1        | 3      | 26.0 | 7.9250  |
| 3 | 1        | 1      | 35.0 | 53.1000 |
| 4 | 0        | 3      | 35.0 | 8.0500  |

In [17]:



```
# To find missing values from the aforementioned columns, you
# need to first call the isnull() method on the titanic_data
# dataframe, and then you need to call the mean() method, as
# shown below.

titanic_data.isnull().mean()

# The output shows that only the age column contains
# missing values. And the ratio of missing values is around 19.86
# percent.
```

Out[17]:

```
survived    0.000000
pclass      0.000000
age         0.198653
fare        0.000000
dtype: float64
```

In [18]:



```
# Let's now find out the median and mean values for all the nonmissing
# values in the age column.

median = titanic_data.age.median()
print(median)
mean = titanic_data.age.mean()
print(mean)

# The age column has a median value of 28 and a mean value of
# 29.6991.
```

```
28.0
29.69911764705882
```

In [19]:

```
# To plot the kernel density plots for the actual age and median
# and mean age, we will add columns to the Pandas dataframe.

import numpy as np
titanic_data['Median_Age'] = titanic_data.age.fillna(median)
titanic_data['Mean_Age'] = titanic_data.age.fillna(mean)
titanic_data['Mean_Age'] = np.round(titanic_data['Mean_Age'], 1)
titanic_data.head(20)

# The above script adds Median_Age and Mean_Age columns
# to the titanic_data dataframe and prints the first 20 records.
# Here is the output of the above script:
```

Out[19]:

|    | survived | pclass | age  | fare    | Median_Age | Mean_Age |
|----|----------|--------|------|---------|------------|----------|
| 0  | 0        | 3      | 22.0 | 7.2500  | 22.0       | 22.0     |
| 1  | 1        | 1      | 38.0 | 71.2833 | 38.0       | 38.0     |
| 2  | 1        | 3      | 26.0 | 7.9250  | 26.0       | 26.0     |
| 3  | 1        | 1      | 35.0 | 53.1000 | 35.0       | 35.0     |
| 4  | 0        | 3      | 35.0 | 8.0500  | 35.0       | 35.0     |
| 5  | 0        | 3      | NaN  | 8.4583  | 28.0       | 29.7     |
| 6  | 0        | 1      | 54.0 | 51.8625 | 54.0       | 54.0     |
| 7  | 0        | 3      | 2.0  | 21.0750 | 2.0        | 2.0      |
| 8  | 1        | 3      | 27.0 | 11.1333 | 27.0       | 27.0     |
| 9  | 1        | 2      | 14.0 | 30.0708 | 14.0       | 14.0     |
| 10 | 1        | 3      | 4.0  | 16.7000 | 4.0        | 4.0      |
| 11 | 1        | 1      | 58.0 | 26.5500 | 58.0       | 58.0     |
| 12 | 0        | 3      | 20.0 | 8.0500  | 20.0       | 20.0     |
| 13 | 0        | 3      | 39.0 | 31.2750 | 39.0       | 39.0     |
| 14 | 0        | 3      | 14.0 | 7.8542  | 14.0       | 14.0     |
| 15 | 1        | 2      | 55.0 | 16.0000 | 55.0       | 55.0     |
| 16 | 0        | 3      | 2.0  | 29.1250 | 2.0        | 2.0      |
| 17 | 1        | 2      | NaN  | 13.0000 | 28.0       | 29.7     |
| 18 | 0        | 3      | 31.0 | 18.0000 | 31.0       | 31.0     |
| 19 | 1        | 3      | NaN  | 7.2250  | 28.0       | 29.7     |

In [20]:

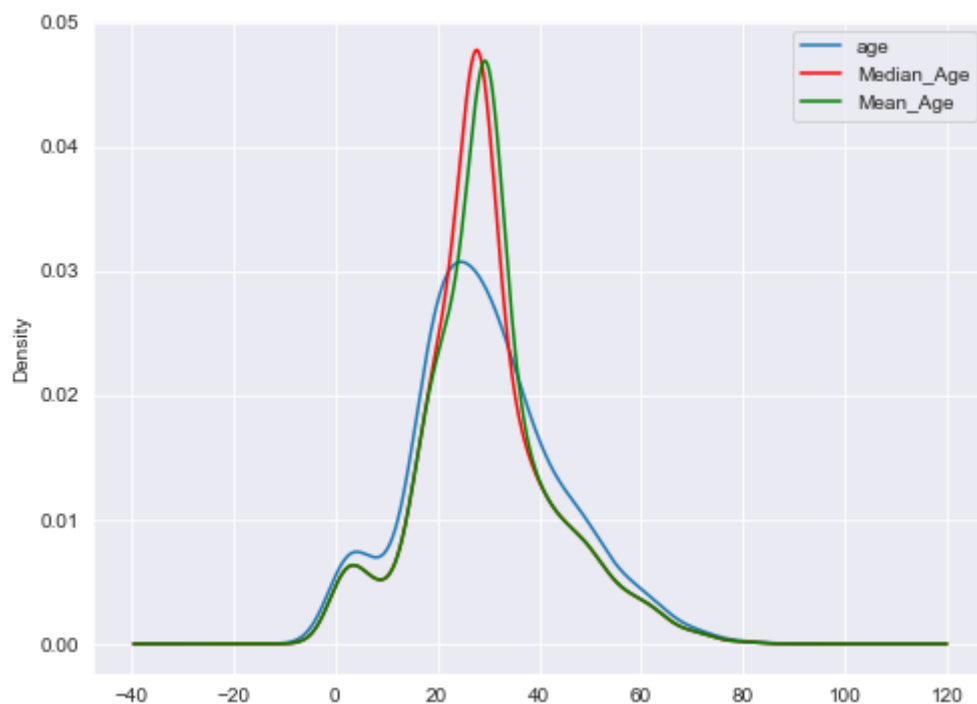
```
# Some rows in the above output show that NaN, i.e.,  
# null values in the age column, have been replaced by the  
# median values in the Median_Age column and by mean values  
# in the Mean_Age column.  
# The mean and median imputation can affect the data  
# distribution for the columns containing the missing values.  
# Specifically, the variance of the column is decreased by mean  
# and median imputation now since more values are added to  
# the center of the distribution. The following script plots the  
# distribution of data for the age, Median_Age, and Mean_Age  
# columns.
```

```
fig = plt.figure()  
ax = fig.add_subplot(111)  
titanic_data['age'].plot(kind='kde', ax=ax)  
titanic_data['Median_Age'].plot(kind='kde', ax=ax, color='red')  
titanic_data['Mean_Age'].plot(kind='kde', ax=ax, color='green')  
lines, labels = ax.get_legend_handles_labels()  
ax.legend(lines, labels, loc='best')
```

```
# Here is the output of the script above:
```

Out[20]:

&lt;matplotlib.legend.Legend at 0x63dfb64a30&gt;



In [ ]:



```
# You can see that the default values in the age columns have  
# been distorted by the mean and median imputation, and the  
# overall variance of the dataset has also been decreased.  
  
#Recommendation  
  
# Mean and Median imputation could be used for the missing  
# numerical data in case the data is missing at random. If the  
# data is normally distributed, mean imputation is better, or else,  
# median imputation is preferred in case of skewed  
# distributions.
```